

Generative AI & Large Language Models

1. Generative Models - Methodology & Data Sources

This below table covers ten LLM releases with confirmed general availability dates from May 2025 onward. Data was cross-referenced from official provider announcements, official API documentation and pricing pages, peer-reviewed benchmark publications, and reputable third-party analysis platforms including Artificial Analysis, LLM Stats, AI Flash Report, and MIT Technology Review. All pricing figures are stated as USD per 1 million tokens (input / output) at standard API rates as of the report compilation date (May 2026), unless otherwise noted. Models with tiered or dynamic pricing reflect standard-tier rates.

The following providers and models are covered in this report:

- OpenAI — GPT-5 (August 2025), GPT-5.5 (April 2026)
- Anthropic — Claude Sonnet 4.6 (February 2026)
- Google DeepMind — Gemini 2.5 Pro (June 2025)
- Meta AI — Llama 4 Scout (April 2025)
- Mistral AI — Mistral Medium 3 (May 2025)
- xAI — Grok 3 Mini (May 2025)
- MiniMax — MiniMax M1 (June 2025)
- Alibaba / Qwen Team — Qwen3 235B-A22B (April 2025)
- DeepSeek AI — DeepSeek V4 (April 2026)

2. Comparative Model Overview

The table below provides a structured, side-by-side comparison of all ten models across eight key analytical dimensions. Pricing is expressed per 1 million tokens (input / output) at standard API rates. Context window figures represent maximum supported input size. Use case and feature descriptions are derived exclusively from official provider documentation and verified third-party benchmark sources.

Provider	Model Name	Release Date	Accessibility	Token Cost (per 1M tokens)	Context Window	Primary Use Cases
OpenAI	GPT-5	August 2025	Public	\$2.50 / \$12.00 per 1M tokens (in/out)	128K tokens	Reasoning, coding, conversational AI, multi-step tool use, agentic workflows
OpenAI	GPT-5.5	April 2026	Public	\$5.00 / \$30.00 per 1M tokens (in/out)	1,050K tokens	Frontier reasoning, agentic coding, knowledge work, computer use, scientific research
Anthropic	Claude Sonnet 4.6	February 2026	Public	\$3.00 / \$15.00 per 1M tokens (in/out)	500K tokens	Agentic workflows, software engineering, long-context document analysis, enterprise coding
Google DeepMind	Gemini 2.5 Pro	June 2025	Public	\$1.25 / \$10.00 per 1M tokens (in/out)	1M tokens	Multimodal reasoning, code generation, long-document analysis, video/audio understanding
Meta AI	Llama 4 Scout	April 2025	Public (Open Weight)	\$0.08 / \$0.30 per 1M tokens (via third-party APIs)	10M tokens	Multimodal document analysis, research, long-context reasoning, on-device deployment
Mistral AI	Mistral Medium 3	May 2025	Public	\$0.40 / \$2.00 per 1M tokens (in/out)	131K tokens	Enterprise text + image tasks, structured output, function calling, cost-efficient production workloads
xAI	Grok 3 Mini	May 2025	Public	Available via X Premium / API (pricing not officially published per 1K tokens)	131K tokens	Cost-efficient reasoning, mathematics, STEM problem solving, real-time information retrieval
MiniMax	MiniMax M1	June 2025	Public (Open Source)	\$0.40 / \$2.20 per 1M tokens (in/out)	1M tokens	Agentic workflows, interactive agents at scale,

Provider	Model Name	Release Date	Accessibility	Token Cost (per 1M tokens)	Context Window	Primary Use Cases
						tool use, cost-sensitive production pipelines
Alibaba / Qwen Team	Qwen3 235B-A22B	April 2025	Public (Open Source)	\$0.10 / \$0.40 per 1M tokens (in/out, via Alibaba Cloud)	128K tokens (extendable to 131K via YaRN)	Reasoning, code generation, multilingual NLP, agentic tasks, local/server MoE deployment
DeepSeek AI	DeepSeek V4	April 2026	Public (Open Source)	\$1.74 / \$3.48 per 1M tokens (V4-Pro); \$0.14 / \$0.28 (V4-Flash)	1M tokens	Agentic coding, STEM reasoning, long-context analysis, cost-efficient frontier-grade pipelines

